

Sentiment Analysis using NLP

March 16, 2025

```
[1]: import pandas as pd
import re
import nltk
from nltk.corpus import stopwords
from nltk.tokenize import word_tokenize
from nltk.stem import PorterStemmer, WordNetLemmatizer
from nltk.tag import pos_tag
from nltk.chunk import ne_chunk

# Load dataset
data = pd.read_csv('kindle_review.csv')
df = data[['reviewText', 'rating']].copy()
```

```
[2]: df
```

```
[2]:
```

	reviewText	rating
0	Jace Rankin may be short, but he's nothing to ...	3
1	Great short read. I didn't want to put it dow...	5
2	I'll start by saying this is the first of four...	3
3	Aggie is Angela Lansbury who carries pocketboo...	3
4	I did not expect this type of book to be in li...	4
...
11995	Valentine cupid is a vampire- Jena and Ian ano...	4
11996	I have read all seven books in this series. Ap...	5
11997	This book really just wasn't my cuppa. The si...	3
11998	tried to use it to charge my kindle, it didn't...	1
11999	Taking Instruction is a look into the often hi...	3

[12000 rows x 2 columns]

```
[3]: # Convert rating to binary (0: negative, 1: positive)
df['rating'] = df['rating'].apply(lambda x: 0 if x < 3 else 1)

# Lowercase
df['reviewText'] = df['reviewText'].str.lower()

# Remove special characters, URLs, and HTML tags
```

```

df['reviewText'] = df['reviewText'].apply(lambda x: re.sub(r'[~a-zA-Z0-9\s]', ' ', str(x)))
df['reviewText'] = df['reviewText'].apply(lambda x: re.sub(r'http\S+', ' ', x))

# Download NLTK resources
nltk.download(['punkt', 'stopwords', 'wordnet', 'averaged_perceptron_tagger', 'maxent_ne_chunker', 'words'])

# Tokenization
df['tokens'] = df['reviewText'].apply(word_tokenize)

# Remove stopwords
stop_words = set(stopwords.words('english'))
df['tokens'] = df['tokens'].apply(lambda tokens: [word for word in tokens if word not in stop_words])

# Stemming
stemmer = PorterStemmer()
df['stemmed'] = df['tokens'].apply(lambda tokens: [stemmer.stem(word) for word in tokens])

# POS Tagging
df['pos_tags'] = df['tokens'].apply(pos_tag)

# Lemmatization with POS
lemmatizer = WordNetLemmatizer()
def lemmatize_with_pos(tagged_tokens):
    lemmatized = []
    for word, tag in tagged_tokens:
        pos = tag[0].lower()
        pos = pos if pos in ['a', 'r', 'n', 'v'] else 'n' # Default to noun
        lemmatized.append(lemmatizer.lemmatize(word, pos))
    return lemmatized

df['lemmatized'] = df['pos_tags'].apply(lemmatize_with_pos)

```

```

[nltk_data] Downloading package punkt to
[nltk_data] C:\Users\itzsh\AppData\Roaming\nltk_data...
[nltk_data] Package punkt is already up-to-date!
[nltk_data] Downloading package stopwords to
[nltk_data] C:\Users\itzsh\AppData\Roaming\nltk_data...
[nltk_data] Package stopwords is already up-to-date!
[nltk_data] Downloading package wordnet to
[nltk_data] C:\Users\itzsh\AppData\Roaming\nltk_data...
[nltk_data] Package wordnet is already up-to-date!
[nltk_data] Downloading package averaged_perceptron_tagger to
[nltk_data] C:\Users\itzsh\AppData\Roaming\nltk_data...

```

```
[nltk_data] Package averaged_perceptron_tagger is already up-to-
[nltk_data] date!
[nltk_data] Downloading package maxent_ne_chunker to
[nltk_data] C:\Users\itzsh\AppData\Roaming\nltk_data...
[nltk_data] Package maxent_ne_chunker is already up-to-date!
[nltk_data] Downloading package words to
[nltk_data] C:\Users\itzsh\AppData\Roaming\nltk_data...
[nltk_data] Package words is already up-to-date!
```

```
[4]: from sklearn.feature_extraction.text import CountVectorizer, TfidfVectorizer
from gensim.models import Word2Vec
import numpy as np

# Bag of Words (BOW)
bow_vectorizer = CountVectorizer()
X_bow = bow_vectorizer.fit_transform(df['lemmatized'].apply(' '.join))

# TF-IDF
tfidf_vectorizer = TfidfVectorizer()
X_tfidf = tfidf_vectorizer.fit_transform(df['lemmatized'].apply(' '.join))

# Word2Vec
# Train Word2Vec model
sentences = df['lemmatized'].tolist()
w2v_model = Word2Vec(sentences, vector_size=100, window=5, min_count=1,
    ↪workers=4)

# Convert sentences to vectors by averaging word vectors
def sentence_vector(sentence):
    return np.mean([w2v_model.wv[word] for word in sentence if word in
    ↪w2v_model.wv], axis=0)

X_w2v = np.array([sentence_vector(sentence) for sentence in sentences])
```

```
[5]: df
```

```
[5]:
```

	reviewText	rating	\
0	jace rankin may be short but hes nothing to me...	1	
1	great short read i didnt want to put it down ...	1	
2	ill start by saying this is the first of four ...	1	
3	aggie is angela lansbury who carries pocketboo...	1	
4	i did not expect this type of book to be in li...	1	
...	
11995	valentine cupid is a vampire jena and ian anot...	1	
11996	i have read all seven books in this series apo...	1	
11997	this book really just wasnt my cuppa the situ...	1	
11998	tried to use it to charge my kindle it didnt e...	0	

11999 taking instruction is a look into the often hi... 1

tokens \

0 [jace, rankin, may, short, hes, nothing, mess,...
1 [great, short, read, didnt, want, put, read, o...
2 [ill, start, saying, first, four, books, wasnt...
3 [aggie, angela, lansbury, carries, pocketbooks...
4 [expect, type, book, library, pleased, find, p...
...
11995 [valentine, cupid, vampire, jena, ian, another...
11996 [read, seven, books, series, apocalypticadvent...
11997 [book, really, wasnt, cuppa, situation, man, c...
11998 [tried, use, charge, kindle, didnt, even, regi...
11999 [taking, instruction, look, often, hidden, wor...

stemmed \

0 [jace, rankin, may, short, he, noth, mess, man...
1 [great, short, read, didnt, want, put, read, o...
2 [ill, start, say, first, four, book, wasnt, ex...
3 [aggi, angela, lansburi, carri, pocketbook, in...
4 [expect, type, book, librari, pleas, find, pri...
...
11995 [valentin, cupid, vampir, jena, ian, anoth, va...
11996 [read, seven, book, seri, apocalypticadventur,...
11997 [book, realli, wasnt, cuppa, situat, man, capt...
11998 [tri, use, charg, kindl, didnt, even, regist, ...
11999 [take, instruct, look, often, hidden, world, s...

pos_tags \

0 [(jace, NN), (rankin, NN), (may, MD), (short, ...
1 [(great, JJ), (short, JJ), (read, NN), (didnt,...
2 [(ill, JJ), (start, VB), (saying, VBG), (first...
3 [(aggie, NN), (angela, JJ), (lansbury, NN), (c...
4 [(expect, VB), (type, NN), (book, NN), (librar...
...
11995 [(valentine, NN), (cupid, NN), (vampire, NN), ...
11996 [(read, VB), (seven, CD), (books, NNS), (serie...
11997 [(book, NN), (really, RB), (wasnt, JJ), (cuppa...
11998 [(tried, VBN), (use, JJ), (charge, NN), (kindl...
11999 [(taking, VBG), (instruction, NN), (look, NN),...

lemmatized

0 [jace, rankin, may, short, he, nothing, mess, ...
1 [great, short, read, didnt, want, put, read, o...
2 [ill, start, say, first, four, book, wasnt, ex...
3 [aggie, angela, lansbury, carry, pocketbook, i...
4 [expect, type, book, library, please, find, pr...

```

...
11995 [valentine, cupid, vampire, jena, ian, another...
11996 [read, seven, book, series, apocalypticadventu...
11997 [book, really, wasnt, cuppa, situation, man, c...
11998 [try, use, charge, kindle, didnt, even, regist...
11999 [take, instruction, look, often, hidden, world...

```

[12000 rows x 6 columns]

```

[5]: from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import accuracy_score
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report

# Split data
X_train_bow, X_test_bow, y_train, y_test = train_test_split(X_bow,
    ↪df['rating'], test_size=0.2)
X_train_tfidf, X_test_tfidf, _, _ = train_test_split(X_tfidf, df['rating'],
    ↪test_size=0.2)
X_train_w2v, X_test_w2v, _, _ = train_test_split(X_w2v, df['rating'],
    ↪test_size=0.2)

# Train models
bow = GaussianNB().fit(X_train_bow.toarray(), y_train)
tfidf = GaussianNB().fit(X_train_tfidf.toarray(), y_train)
w2v = GaussianNB().fit(X_train_w2v, y_train)

# Evaluate
y_pred_bow = bow.predict(X_test_bow.toarray())
y_pred_tfidf = tfidf.predict(X_test_tfidf.toarray())
y_pred_w2v = w2v.predict(X_test_w2v)

print("BOW Accuracy:", accuracy_score(y_test, y_pred_bow))
print("TF-IDF Accuracy:", accuracy_score(y_test, y_pred_tfidf))
print("Word2Vec Accuracy:", accuracy_score(y_test, y_pred_w2v))

```

BOW Accuracy: 0.58625

TF-IDF Accuracy: 0.4483333333333333

Word2Vec Accuracy: 0.55875

```

[6]: # Train models
bow = RandomForestClassifier().fit(X_train_bow.toarray(), y_train)
tfidf = RandomForestClassifier().fit(X_train_tfidf.toarray(), y_train)
w2v = RandomForestClassifier().fit(X_train_w2v, y_train)

```

```

# Evaluate
y_pred_bow = bow.predict(X_test_bow.toarray())
y_pred_tfidf = tfidf.predict(X_test_tfidf.toarray())
y_pred_w2v = w2v.predict(X_test_w2v)

print("BOW Accuracy:", accuracy_score(y_test, y_pred_bow))
print("TF-IDF Accuracy:", accuracy_score(y_test, y_pred_tfidf))
print("Word2Vec Accuracy:", accuracy_score(y_test, y_pred_w2v))

```

BOW Accuracy: 0.795
 TF-IDF Accuracy: 0.6695833333333333
 Word2Vec Accuracy: 0.665

```

[7]: # Train models
bow = SVC(kernel='linear').fit(X_train_bow.toarray(), y_train)
tfidf = SVC(kernel='linear').fit(X_train_tfidf.toarray(), y_train)
w2v = SVC(kernel='linear').fit(X_train_w2v, y_train)

# Evaluate
y_pred_bow = bow.predict(X_test_bow.toarray())
y_pred_tfidf = tfidf.predict(X_test_tfidf.toarray())
y_pred_w2v = w2v.predict(X_test_w2v)

print("BOW Accuracy:", accuracy_score(y_test, y_pred_bow))
print("TF-IDF Accuracy:", accuracy_score(y_test, y_pred_tfidf))
print("Word2Vec Accuracy:", accuracy_score(y_test, y_pred_w2v))

```

BOW Accuracy: 0.82
 TF-IDF Accuracy: 0.6670833333333334
 Word2Vec Accuracy: 0.675

```

[8]: # TF-IDF Features (reuse existing)
X = tfidf_vectorizer.fit_transform(df['lemmatized'].apply(' '.join))
y = df['rating']

# Split data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    random_state=42)

# SVM Model
svm = SVC(kernel='linear')
svm.fit(X_train, y_train)
y_pred_svm = svm.predict(X_test)
print("SVM Report:\n", classification_report(y_test, y_pred_svm))
print(accuracy_score(y_test, y_pred_svm))

```

SVM Report:

precision	recall	f1-score	support
-----------	--------	----------	---------

0	0.82	0.70	0.75	803
1	0.86	0.92	0.89	1597
accuracy				0.85 2400
macro avg	0.84	0.81	0.82	2400
weighted avg	0.85	0.85	0.84	2400

0.8470833333333333

```
[9]: # Random Forest Model
rf = RandomForestClassifier(n_estimators=100, class_weight='balanced')
rf.fit(X_train, y_train)
y_pred_rf = rf.predict(X_test)
print("Random Forest Report:\n", classification_report(y_test, y_pred_rf))
print(accuracy_score(y_test, y_pred_rf))
```

Random Forest Report:

	precision	recall	f1-score	support
0	0.86	0.45	0.60	803
1	0.78	0.96	0.86	1597
accuracy				0.79 2400
macro avg	0.82	0.71	0.73	2400
weighted avg	0.81	0.79	0.77	2400

0.7933333333333333

```
[10]: # TF-IDF Features (reuse existing)
X = bow_vectorizer.fit_transform(df['lemmatized'].apply(' '.join))
y = df['rating']

# Split data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳ random_state=42)

# SVM Model
svm = SVC(kernel='linear')
svm.fit(X_train, y_train)
y_pred_svm = svm.predict(X_test)
print("SVM Report:\n", classification_report(y_test, y_pred_svm))
print(accuracy_score(y_test, y_pred_svm))
```

SVM Report:

	precision	recall	f1-score	support
0	0.74	0.71	0.73	803
1	0.86	0.87	0.86	1597

accuracy			0.82	2400
macro avg	0.80	0.79	0.79	2400
weighted avg	0.82	0.82	0.82	2400

0.81875

```
[11]: # Random Forest Model
rf = RandomForestClassifier(n_estimators=100, class_weight='balanced')
rf.fit(X_train, y_train)
y_pred_rf = rf.predict(X_test)
print("Random Forest Report:\n", classification_report(y_test, y_pred_rf))
print(accuracy_score(y_test, y_pred_rf))
```

Random Forest Report:

	precision	recall	f1-score	support
0	0.85	0.46	0.59	803
1	0.78	0.96	0.86	1597

accuracy			0.79	2400
macro avg	0.81	0.71	0.73	2400
weighted avg	0.80	0.79	0.77	2400

0.79125

```
[12]: # CNN Model using existing Word2Vec embeddings
import numpy as np
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Embedding, Conv1D, GlobalMaxPooling1D, Dense
from sklearn.model_selection import train_test_split

# Convert lemmatized tokens to sequences
tokenizer = Tokenizer()
tokenizer.fit_on_texts(df['lemmatized'].apply(lambda x: ' '.join(x)))
sequences = tokenizer.texts_to_sequences(df['lemmatized'].apply(lambda x: ' '.
    ↪join(x)))

# Pad sequences
max_len = max(len(s) for s in sequences)
X_cnn = pad_sequences(sequences, maxlen=max_len, padding='post')
y = df['rating']

# Split data with fixed random_state
X_train_cnn, X_test_cnn, y_train_cnn, y_test_cnn = train_test_split(X_cnn, y,
    ↪test_size=0.2, random_state=42)
```



```

# Load existing Word2Vec embeddings
embedding_dim = 100
vocab_size = len(tokenizer.word_index) + 1
embedding_matrix = np.zeros((vocab_size, embedding_dim))
for word, i in tokenizer.word_index.items():
    if word in w2v_model.wv:
        embedding_matrix[i] = w2v_model.wv[word]

# Build CNN model
model_cnn = Sequential()
model_cnn.add(Embedding(vocab_size, embedding_dim, weights=[embedding_matrix],
    ↪input_length=max_len, trainable=False))
model_cnn.add(Conv1D(128, 5, activation='relu'))
model_cnn.add(GlobalMaxPooling1D())
model_cnn.add(Dense(1, activation='sigmoid'))
model_cnn.compile(optimizer='adam', loss='binary_crossentropy',
    ↪metrics=['accuracy'])

# Train and evaluate
model_cnn.fit(X_train_cnn, y_train_cnn, epochs=10, validation_data=(X_test_cnn,
    ↪y_test_cnn))
loss, accuracy = model_cnn.evaluate(X_test_cnn, y_test_cnn)
print(f"CNN Accuracy: {accuracy:.4f}")

```

C:\Users\itzsh\AppData\Local\Programs\Python\Python312\Lib\site-packages\keras\src\layers\core\embedding.py:90: UserWarning: Argument `input_length` is deprecated. Just remove it.
 warnings.warn(

```

Epoch 1/10
300/300          22s 71ms/step -
accuracy: 0.6899 - loss: 0.5924 - val_accuracy: 0.7575 - val_loss: 0.4763
Epoch 2/10
300/300          21s 68ms/step -
accuracy: 0.7940 - loss: 0.4433 - val_accuracy: 0.7254 - val_loss: 0.5694
Epoch 3/10
300/300          20s 68ms/step -
accuracy: 0.7972 - loss: 0.4316 - val_accuracy: 0.7879 - val_loss: 0.4498
Epoch 4/10
300/300          20s 68ms/step -
accuracy: 0.8234 - loss: 0.3829 - val_accuracy: 0.7821 - val_loss: 0.4435
Epoch 5/10
300/300          21s 68ms/step -
accuracy: 0.8467 - loss: 0.3513 - val_accuracy: 0.7812 - val_loss: 0.4414
Epoch 6/10
300/300          20s 68ms/step -
accuracy: 0.8551 - loss: 0.3313 - val_accuracy: 0.7892 - val_loss: 0.4376
Epoch 7/10

```

```

300/300          20s 68ms/step -
accuracy: 0.8788 - loss: 0.2933 - val_accuracy: 0.7987 - val_loss: 0.4308
Epoch 8/10
300/300          20s 68ms/step -
accuracy: 0.8835 - loss: 0.2862 - val_accuracy: 0.7817 - val_loss: 0.4555
Epoch 9/10
300/300          21s 69ms/step -
accuracy: 0.8995 - loss: 0.2552 - val_accuracy: 0.7550 - val_loss: 0.6182
Epoch 10/10
300/300          21s 69ms/step -
accuracy: 0.9100 - loss: 0.2302 - val_accuracy: 0.7829 - val_loss: 0.4664
75/75           1s 18ms/step -
accuracy: 0.7766 - loss: 0.4621
CNN Accuracy: 0.7829

```

```

[1]: import numpy as np
import pandas as pd
import tensorflow as tf
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Embedding, Conv1D, GlobalMaxPooling1D, Dense, Dropout, BatchNormalization, Bidirectional, LSTM, SpatialDropout1D
from tensorflow.keras.regularizers import l2
from tensorflow.keras.optimizers.schedules import ExponentialDecay
from sklearn.model_selection import train_test_split
from gensim.models import Word2Vec

# Load dataset
df = pd.read_csv('kindle_review.csv')
df['rating'] = df['rating'].apply(lambda x: 0 if x < 3 else 1)

# Tokenization
tokenizer = Tokenizer()
tokenizer.fit_on_texts(df['reviewText'])
sequences = tokenizer.texts_to_sequences(df['reviewText'])
vocab_size = len(tokenizer.word_index) + 1

# Padding sequences
max_len = 200
X = pad_sequences(sequences, maxlen=max_len, padding='post')
y = df['rating']

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

```

```

# Train Word2Vec model
sentences = [text.split() for text in df['reviewText']]
w2v_model = Word2Vec(sentences, vector_size=100, window=5, min_count=1,
    ↪workers=4)

# Create an embedding matrix
embedding_dim = 100
embedding_matrix = np.zeros((vocab_size, embedding_dim))
for word, i in tokenizer.word_index.items():
    if word in w2v_model.wv:
        embedding_matrix[i] = w2v_model.wv[word]

# Learning rate scheduling
lr_schedule = ExponentialDecay(initial_learning_rate=0.001, decay_steps=10000,
    ↪decay_rate=0.9, staircase=True)

# Build CNN + BiLSTM model
model = Sequential([
    Embedding(vocab_size, embedding_dim, weights=[embedding_matrix],
    ↪input_length=max_len, trainable=True),
    SpatialDropout1D(0.3),

    # CNN Layers
    Conv1D(128, 5, activation='relu', padding='same'),
    BatchNormalization(),
    Dropout(0.3),

    Conv1D(64, 3, activation='relu', padding='same'),
    BatchNormalization(),
    Dropout(0.3),

    # Bidirectional LSTM for contextual understanding
    Bidirectional(LSTM(64, return_sequences=True, dropout=0.3,
    ↪recurrent_dropout=0.3)),
    GlobalMaxPooling1D(),

    # Fully Connected Layers
    Dense(128, activation='relu', kernel_regularizer=l2(0.01)),
    Dropout(0.3),

    Dense(1, activation='sigmoid')
])

# Compile model
model.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=lr_schedule),
    loss='binary_crossentropy',
    metrics=['accuracy'])

```

```

# Train model
model.fit(X_train, y_train, epochs=10, batch_size=64, validation_data=(X_test,
↪y_test))

# Evaluate model
loss, accuracy = model.evaluate(X_test, y_test)
print(f"Improved CNN + LSTM Accuracy with Word2Vec: {accuracy:.4f}")

```

Epoch 1/10

```

C:\Users\itzsh\AppData\Local\Programs\Python\Python312\Lib\site-
packages\keras\src\layers\core\embedding.py:90: UserWarning: Argument
`input_length` is deprecated. Just remove it.
  warnings.warn(

```

```

150/150          80s 497ms/step -
accuracy: 0.6526 - loss: 1.4850 - val_accuracy: 0.7096 - val_loss: 0.7683

```

Epoch 2/10

```

150/150          80s 532ms/step -
accuracy: 0.7241 - loss: 0.6896 - val_accuracy: 0.6792 - val_loss: 0.6379

```

Epoch 3/10

```

150/150          82s 548ms/step -
accuracy: 0.7747 - loss: 0.5280 - val_accuracy: 0.8179 - val_loss: 0.4324

```

Epoch 4/10

```

150/150          81s 543ms/step -
accuracy: 0.8103 - loss: 0.4605 - val_accuracy: 0.8158 - val_loss: 0.4199

```

Epoch 5/10

```

150/150          81s 539ms/step -
accuracy: 0.8221 - loss: 0.4156 - val_accuracy: 0.8071 - val_loss: 0.4201

```

Epoch 6/10

```

150/150          82s 546ms/step -
accuracy: 0.8444 - loss: 0.3751 - val_accuracy: 0.8117 - val_loss: 0.4123

```

Epoch 7/10

```

150/150          82s 545ms/step -
accuracy: 0.8604 - loss: 0.3383 - val_accuracy: 0.8575 - val_loss: 0.3531

```

Epoch 8/10

```

150/150          83s 554ms/step -
accuracy: 0.8867 - loss: 0.2950 - val_accuracy: 0.8504 - val_loss: 0.3588

```

Epoch 9/10

```

150/150          81s 544ms/step -
accuracy: 0.9091 - loss: 0.2481 - val_accuracy: 0.8367 - val_loss: 0.3840

```

Epoch 10/10

```

150/150          82s 545ms/step -
accuracy: 0.9290 - loss: 0.2047 - val_accuracy: 0.8400 - val_loss: 0.4293

```

```

75/75           2s 32ms/step -
accuracy: 0.8322 - loss: 0.4397

```

Improved CNN + LSTM Accuracy with Word2Vec: 0.8400

```
[1]: import numpy as np
import pandas as pd
import tensorflow as tf
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Embedding, Conv1D, GlobalMaxPooling1D,
↳Dense, Dropout, BatchNormalization, Bidirectional, LSTM, SpatialDropout1D
from tensorflow.keras.regularizers import l2
from tensorflow.keras.optimizers.schedules import ExponentialDecay
from sklearn.model_selection import train_test_split
import gensim.downloader as api

# Load dataset
df = pd.read_csv('kindle_review.csv')
df['rating'] = df['rating'].apply(lambda x: 0 if x < 3 else 1)

# Tokenization
tokenizer = Tokenizer()
tokenizer.fit_on_texts(df['reviewText'])
sequences = tokenizer.texts_to_sequences(df['reviewText'])
vocab_size = len(tokenizer.word_index) + 1

# Padding sequences
max_len = 200
X = pad_sequences(sequences, maxlen=max_len, padding='post')
y = df['rating']

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳random_state=42)

# Load pre-trained Glove embeddings
glove_model = api.load("glove-wiki-gigaword-100")
embedding_dim = 100
embedding_matrix = np.zeros((vocab_size, embedding_dim))
for word, i in tokenizer.word_index.items():
    if word in glove_model:
        embedding_matrix[i] = glove_model[word]

# Learning rate scheduling
lr_schedule = ExponentialDecay(initial_learning_rate=0.001, decay_steps=10000,
↳decay_rate=0.9, staircase=True)

# Build improved CNN + LSTM model
model = Sequential([
```

```

    Embedding(vocab_size, embedding_dim, weights=[embedding_matrix],  

    ↪input_length=max_len, trainable=True),  

    SpatialDropout1D(0.3),  
  

    # First Convolutional Block  

    Conv1D(128, 5, activation='relu', padding='same'),  

    BatchNormalization(),  

    Dropout(0.3),  
  

    # Second Convolutional Block  

    Conv1D(64, 3, activation='relu', padding='same'),  

    BatchNormalization(),  

    Dropout(0.3),  
  

    # LSTM Layer with Residual Connection  

    Bidirectional(LSTM(64, return_sequences=True, dropout=0.3,  

    ↪recurrent_dropout=0.3)),  

    GlobalMaxPooling1D(),  
  

    Dense(128, activation='relu', kernel_regularizer=l2(0.01)),  

    Dropout(0.3),  
  

    Dense(1, activation='sigmoid')  

])  
  

# Compile model  

model.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=lr_schedule),  

              loss='binary_crossentropy',  

              metrics=['accuracy'])  
  

# Train model  

model.fit(X_train, y_train, epochs=10, batch_size=64, validation_data=(X_test,  

    ↪y_test))  
  

# Evaluate model  

loss, accuracy = model.evaluate(X_test, y_test)  

print(f"Improved CNN + LSTM Accuracy: {accuracy:.4f}")

```

C:\Users\itzsh\AppData\Local\Programs\Python\Python312\Lib\site-packages\keras\src\layers\core\embedding.py:90: UserWarning: Argument `input_length` is deprecated. Just remove it.
 warnings.warn(

Epoch 1/10

150/150 89s 558ms/step -

accuracy: 0.6553 - loss: 1.4786 - val_accuracy: 0.7058 - val_loss: 0.7182

Epoch 2/10

150/150 83s 553ms/step -

```

accuracy: 0.7198 - loss: 0.6739 - val_accuracy: 0.8129 - val_loss: 0.4527
Epoch 3/10
150/150      89s 594ms/step -
accuracy: 0.8028 - loss: 0.4815 - val_accuracy: 0.8392 - val_loss: 0.3916
Epoch 4/10
150/150      84s 558ms/step -
accuracy: 0.8342 - loss: 0.4070 - val_accuracy: 0.8254 - val_loss: 0.4377
Epoch 5/10
150/150      84s 557ms/step -
accuracy: 0.8586 - loss: 0.3529 - val_accuracy: 0.8250 - val_loss: 0.3922
Epoch 6/10
150/150      84s 560ms/step -
accuracy: 0.8825 - loss: 0.2988 - val_accuracy: 0.8621 - val_loss: 0.3392
Epoch 7/10
150/150      84s 562ms/step -
accuracy: 0.8994 - loss: 0.2620 - val_accuracy: 0.8596 - val_loss: 0.3382
Epoch 8/10
150/150      83s 554ms/step -
accuracy: 0.9179 - loss: 0.2352 - val_accuracy: 0.8446 - val_loss: 0.4523
Epoch 9/10
150/150      86s 574ms/step -
accuracy: 0.9265 - loss: 0.2077 - val_accuracy: 0.8604 - val_loss: 0.3563
Epoch 10/10
150/150      93s 624ms/step -
accuracy: 0.9286 - loss: 0.2008 - val_accuracy: 0.8596 - val_loss: 0.3646
75/75        3s 38ms/step -
accuracy: 0.8515 - loss: 0.3746
Improved CNN + LSTM Accuracy: 0.8596

```